

## פרק 5

# חולשות (המשק)

## Heap / זיכרון



# Integer overflow

• מה יודפס?

```
4  int _tmain(int argc, _TCHAR* argv[])
5  {
6      int a=-5;
7      unsigned int b=80;
8
9      if ( (unsigned int )a < b )
10         printf(" %d < %d\n",a,b);
11     else
12         printf(" %d >= %d\n",a,b);
13
14     return 0;
15 }
16
```



# Stagefright (2015)

CVE-ID	
<b>CVE-2015-3864</b>	<a href="#">Learn more at National Vulnerability Database (NVD)</a> • Severity Rating • Fix Information • Vulnerable Software Versions • SCAP Mappings
Description	
Integer underflow in the MPEG4Extractor::parseChunk function in MPEG4Extractor.cpp in libstagefright in mediaserver in Android before 5.1.1 LMY48M allows remote attackers to execute arbitrary code via crafted MPEG-4 data, aka internal bug 23034759. NOTE: this vulnerability exists because of an incomplete fix for CVE-2015-3824.	
References	
<b>Note:</b> <a href="#">References</a> are provided for the convenience of the reader to help distinguish between vulnerabilities. The list is not intended to be complete.	
<ul style="list-style-type: none"> <li>• MLIST:[android-security-updates] 20150909 Nexus Security Bulletin (September 2015)</li> <li>• URL:<a href="https://groups.google.com/forum/message/raw?msg=android-security-updates/1M7qbSvACjo/Y7jewiW1AwAJ">https://groups.google.com/forum/message/raw?msg=android-security-updates/1M7qbSvACjo/Y7jewiW1AwAJ</a></li> <li>• MISC:<a href="https://blog.zimperium.com/cve-2015-3864-metasploit-module-now-available-for-testing/">https://blog.zimperium.com/cve-2015-3864-metasploit-module-now-available-for-testing/</a></li> <li>• MISC:<a href="https://blog.zimperium.com/reflecting-on-stagefright-patches/">https://blog.zimperium.com/reflecting-on-stagefright-patches/</a></li> <li>• CONFIRM:<a href="https://android.googlesource.com/platform/frameworks/av/+/-/6fe85f7e15203e48df2cc3e8e1c4bc6ad49dc968">https://android.googlesource.com/platform/frameworks/av/+/-/6fe85f7e15203e48df2cc3e8e1c4bc6ad49dc968</a></li> <li>• BID:76682</li> <li>• URL:<a href="https://www.securityfocus.com/bid/76682">https://www.securityfocus.com/bid/76682</a></li> </ul>	

## Stagefright (bug)

From Wikipedia, the free encyclopedia

**Stagefright** is the group of [software bugs](#) that affect versions 2.2 ("Froyo") and newer of the [Android operating system](#), allowing an attacker to perform arbitrary operations on the victim's device through [remote code execution](#) and [privilege escalation](#).<sup>[1]</sup> Security researchers demonstrate the bugs with a [proof of concept](#) that sends specially crafted MMS messages to the victim device and in most cases requires no [end-user](#) actions upon message reception to succeed - the user doesn't have to do anything to 'accept' the bug - it happens in the background. The phone number is the only target information.<sup>[2][3][4][5]</sup>

The underlying [attack vector](#) exploits certain [integer overflow vulnerabilities](#) in the Android core component called "Stagefright",<sup>[6][7][a]</sup> which is a complex [software library](#) implemented primarily in C++ as part of the Android Open Source Project (AOSP) and used as a backend engine for playing various multimedia formats such as MP4 files.<sup>[5][9]</sup>

The discovered bugs have been provided with multiple [Common Vulnerabilities and Exposures \(CVE\) identifiers](#), CVE-2015-1538, CVE-2015-1539, CVE-2015-3824, CVE-2015-3826, CVE-2015-3827, CVE-2015-3828, CVE-2015-3829 and CVE-2015-3864 (the latter one has been assigned separately from the others), which are collectively referred to as the Stagefright bug.<sup>[10][11][12]</sup>




# דואל אהחייט – Network card driver remote vulnerability

- מצאו את החולשה ב-vuln.c

```
3.
4. /* Snippet code from RT73 USB Network card Driver (sanity.c - version < V1.0.5.0 ) . */
5.
6. #define MAX_LEN_OF_SSID 32
7.
8. typedef struct PACKED _FRAME_802_11 {
9.     HEADER_802_11 Hdr;
10.    CHAR Octet[1];
11. } FRAME_802_11, *PFRAME_802_11;
12.
13. BOOLEAN PeerProbeReqSanity( IN PRTMP_ADAPTER pAd, IN VOID *Msg,
14.    IN ULONG MsgLen, OUT PCHAR pAddr2, OUT CHAR Ssid[], OUT UCHAR *pSsidLen)
15. {
16.    UCHAR Idx;
17.    UCHAR RateLen;
18.    CHAR IEType;
19.    PFRAME_802_11 pFrame = (PFRAME_802_11)Msg;
20.
21.    if ((pFrame->Octet[0] != IE_SSID) || (pFrame->Octet[1] > MAX_LEN_OF_SSID))
22.    {
23.        DBGPRINT(RT_DEBUG_TRACE, "PeerProbeReqSanity fail - wrong SSID IE(Type=%d,Len=%d)\n",
24.            pFrame->Octet[0],pFrame->Octet[1]);
25.        return FALSE;
26.    }
27.
28.    *pSsidLen = pFrame->Octet[1];
29.    memcpy(Ssid, &pFrame->Octet[2], *pSsidLen);
30.    .
31.    .
32. }
33.
```



# Back To Overflow...



## מה הפרימיטיב האצניין?

- לא כל החולשות מחייבות מצב של Buffer overflow
- ישנן חולשות המאפשרות כתיבת מידע למקומות בזיכרון, שתכליתן להשיג מה שמכונה **Write What Where**.
  - לעיתים זה עדיף מ-Buffer Overflow, בעיקר בהיבטים של עקיפת הגנות.
- בדרך כלל מדובר במספר קטן של בתים (4 בתים – למה?), אבל זה בהחלט מספיק כדי שהחולשה תאפשר השתלטות מלאה.
  - ניתן לשכתב RET, Function pointer ועוד...



# מה אפשר לעשות עם WWW?

- כללי המשחק:

- קיימת חולשת WWW
- קיים קוד של התוקף במערכת.
- התוקף רץ בהרשאות נמוכות ורוצה להריץ את הקוד בהרשאות גבוהות (PE)
  - למה?
  - תנו דוגמא לתרחיש עולם אמיתי...

- האם WWW יעזור לנו?

- מי צריך לבצע את הכתיבה באופן אידיאלי?
- מה עוד צריך להתקיים?







- נניח שהיתה פונקציה שניתן לגרום ל-Kernel להריץ באמעות קריאה (לא מתועדת) מ-user land... (CVE-2014-4113)

```
001588E lea     edx, [rsi+10h] ; Length
0015891 lea     r8d, [rsi+8]   ; Alignment
0015895 call    cs: __imp_ProbeForRead
0015898 mov     rbx, [r12]
001589F mov     rdi, [r12+8]
00158A4 mov     rdx, r12
00158A7 lea     rcx, aUserwritewha_1 ; "[+] UserWriteWhatWhere: 0x%p\n"
00158AE call    DbgPrint_0
00158B3 lea     edx, [rsi+10h]
00158B6 lea     rcx, aWrite_what_whe ; "[+] WRITE_WHAT_WHERE Size: 0x%X\n"
00158BD call    DbgPrint_0
00158C2 mov     rdx, rbx
00158C5 lea     rcx, aUserwritewhatw ; "[+] UserWriteWhatWhere->What: 0x%p\n"
00158CC call    DbgPrint_0
00158D1 mov     rdx, rdi
00158D4 lea     rcx, aUserwritewha_0 ; "[+] UserWriteWhatWhere->Where: 0x%p\n"
00158DB call    DbgPrint_0
00158E0 lea     rcx, aTriggeringArbi ; "[+] Triggering Arbitrary Overwrite\n"
00158E7 call    DbgPrint_0
00158EC mov     r11d, [rbx]
00158EF mov     [rdi], r11d
00158F2 jmp     short loc_15C05
```

- R12 נמצא בשליטת התוקף...

- איך אפשר להשתמש בזה?





# מה אפשר לעשות עם זה?

- כתיבת מחדש של Function Pointer
  - איזה Pointer?

- שינוי הרשאות
  - ACL

- שינוי token של התהליך
  - להיות SYSTEM !

- לדרוס כתובת חזרה (בלי Bof !)

כל זאת בתנאי שיודעים איפה רוצים לכתוב...



# חולשות Heap

- Stack Overflow הוא בעיה מוכרת.
- כיצד ניתן להתמודד?
  - Cookies
  - DEP
  - ASLR
  - להעביר את כל ה-buffer-ים ל-Heap
- ניצול חולשת Heap Overflow הוא קשה יותר, אבל אפשרי:

```
void func(void* arg, int len) {  
    char* ptr = malloc(100);  
    memcpy(ptr, arg, len); //buffer overflow if len>100  
    ...  
}
```



# Heap overflow

- חולשות Heap הרבה פחות סטנדרטיות :
  - תלויות בארכיטקטורה
  - המבנים יותר סבוכים, ותלויים במערכת ההפעלה ולעיתים גם בקומפיילר
- המטרה הבסיסית : להגיע מכתובה לא מורשית (Data Overflow) להרצת קוד עוין.



## Use after free

- החולשה : נעשה שימוש במצביע לאחר שהוא שוחרר.
- בדרך-כלל אלו מקרי קצה או צירוף של תנאים שלא היה אמור לקרות, שכיח מאוד בהקשרים של multi-threaded.
- על-ידי מניפולציות על התוכנה, התוקף ינסה לגרום למצביע ששוחרר להצביע על מידע שהוא יכול לשלוט בו.



# Use-After-Free in the Real World

[ThreatPost, September 17, 2013]

The attacks are targeting IE 8 and 9 and there's no patch for the vulnerability right now... **The vulnerability exists in the way that Internet Explorer accesses an object in memory that has been deleted or has not been properly allocated.** The vulnerability may corrupt memory in a way that could allow an attacker to execute arbitrary code...

The exploit was attacking a **Use After Free vulnerability** in IE's HTML rendering engine (mshtml.dll) and was implemented entirely in Javascript (no dependencies on Java, Flash etc), but did depend on a Microsoft Office DLL which was not compiled with ASLR (Address Space Layout Randomization) enabled.

The purpose of this DLL in the context of this exploit is to bypass ASLR by providing executable code at known addresses in memory, so that a hardcoded **ROP (Return Oriented Programming)** chain can be used to mark the pages containing shellcode (in the form of Javascript strings) as executable...

The most likely attack scenarios for this vulnerability are the typical link in an email or drive-by download.

**MICROSOFT WARNS OF NEW IE ZERO DAY, EXPLOIT IN THE WILD**



## Dangling pointer :אנדרג

- מקרים שבהם משתנה האמור להכיל ערכים לשימוש פנימי נשלט על-ידי התוקף.
- נובע בדרך-כלל מטעות של המתכנת.

```
44 void function()  
45 {  
46     list_node current_node;  
47     current_node.id = 1234;  
48     list_insert(list_head, current_node);  
49  
50     .  
51     .  
52     .  
53  
54     if ( problem )  
55     {  
56         printf( "rare error... " );  
57         return;  
58     }  
59  
60     .  
61     .  
62  
63     list_remove(list_head, current_node);  
64  
65 }
```

```
35  
36 typedef struct list_node{  
37     char *data;  
38     int id;  
39     struct list_node *next;  
40 };  
41
```

After function() returns , **current\_node** is dangling on the stack, if the stack will grow enough - **current\_node** will get overwritten.

Linux futex() local vulnerability.



# Double-Free

- שחרור אותו מקטע זיכרון פעמיים.
  - קורה לרוב במקרים של Multi-Threaded
  - ובמקרי קצה...
- למה זה מסוכן?



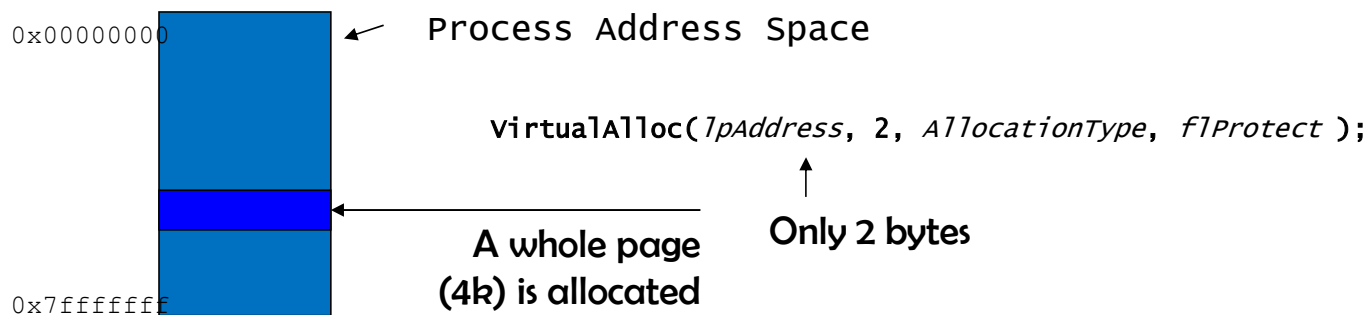


# Digging Deeper...



# The Heap API

The Virtual Memory Manager allocates and deallocates with the granularity of one page (4k)



Even the stack region grows and shrinks in pages (done automatically by the OS)

# The Heap API

- We need a more flexible structure!! We want not to waste space and have allocations of few bytes

```
char* buff = malloc(2 * sizeof(char));
```

and of MBs

```
void* buff = malloc(1048576);
```

- This is done using the structure called *Heap*
- API for the heap is provided by the C runtime library (malloc, free, delete, new)
- In Win32, Heap API is provided by the OS too!



# Heap API

- In Windows mainly following APIs are used to operate on heap.
  - HeapCreate Create a heap
    - \_HEAP structure resides at the beginning of a heap and the address is returned as a HANDLE.
    - \_HEAP structure contains information to manage the heap.
  - HeapAlloc Obtain specific size of memory from heap region
    - Applications can save data into the returned memory region.
    - To manage heap there is management information (Chunk header) just before allocated memory.
    - Regions not allocated yet are managed as free chunks
  - HeapFree Release obtained memory region



# The Heap structures

- In order to handle different size, and cope well with problems (fragmentation) the Heap has a complex structure
- It uses several algorithms too:
  - Allocation algorithm (different strategies for different size)
  - Free algorithm
  - Coalesce (fusion of 2 adjacent free segments)



# The Heap structures

This is the structure that holds Process properties. Its location is fixed at 7ffdf000


(or go via TEB -> PEB)

```
mov eax, dword ptr fs:[18]  
mov eax, dword ptr[eax+0x30]
```

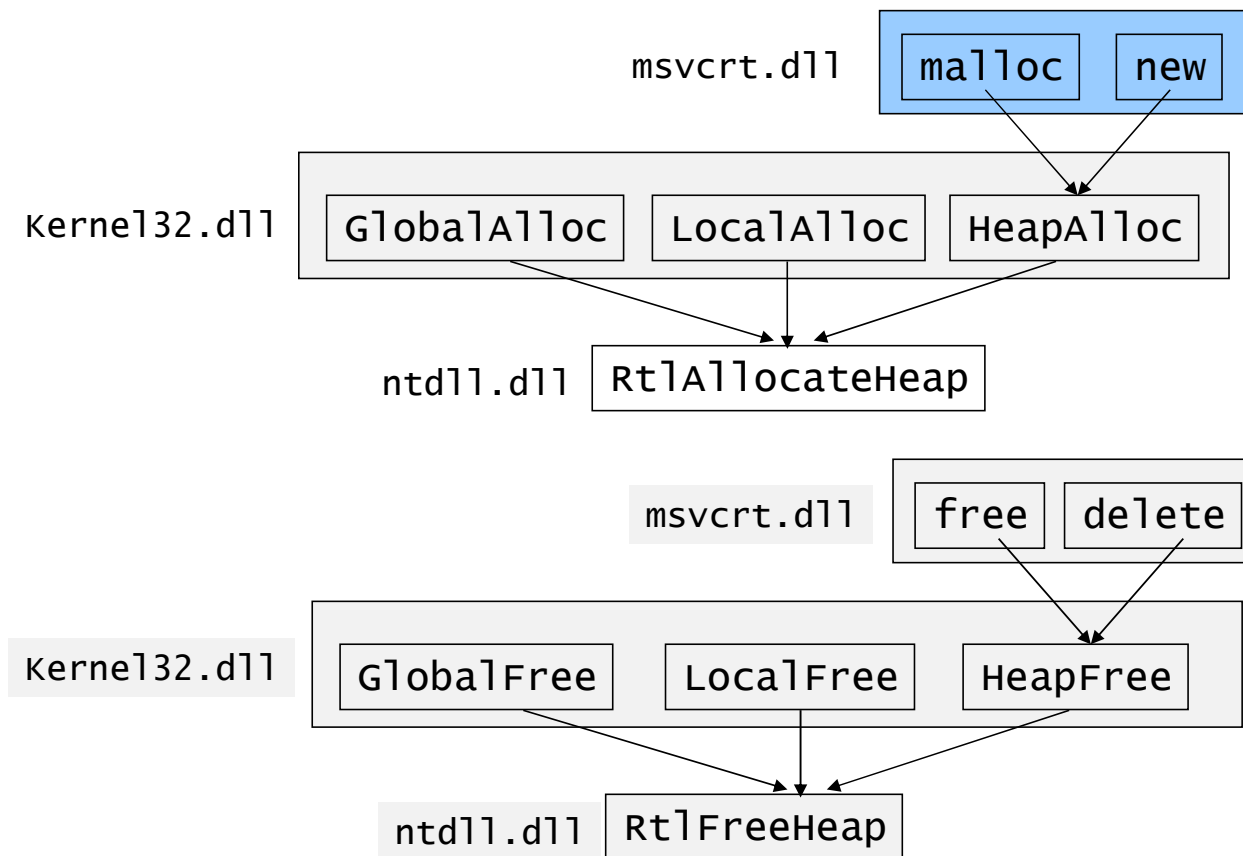
↓  
PEB

0x0010			Default Heap	
0x0080			Heaps Count	
0x0090	Heap List			

So we can easily access to the Heap address.  
This will be useful later!



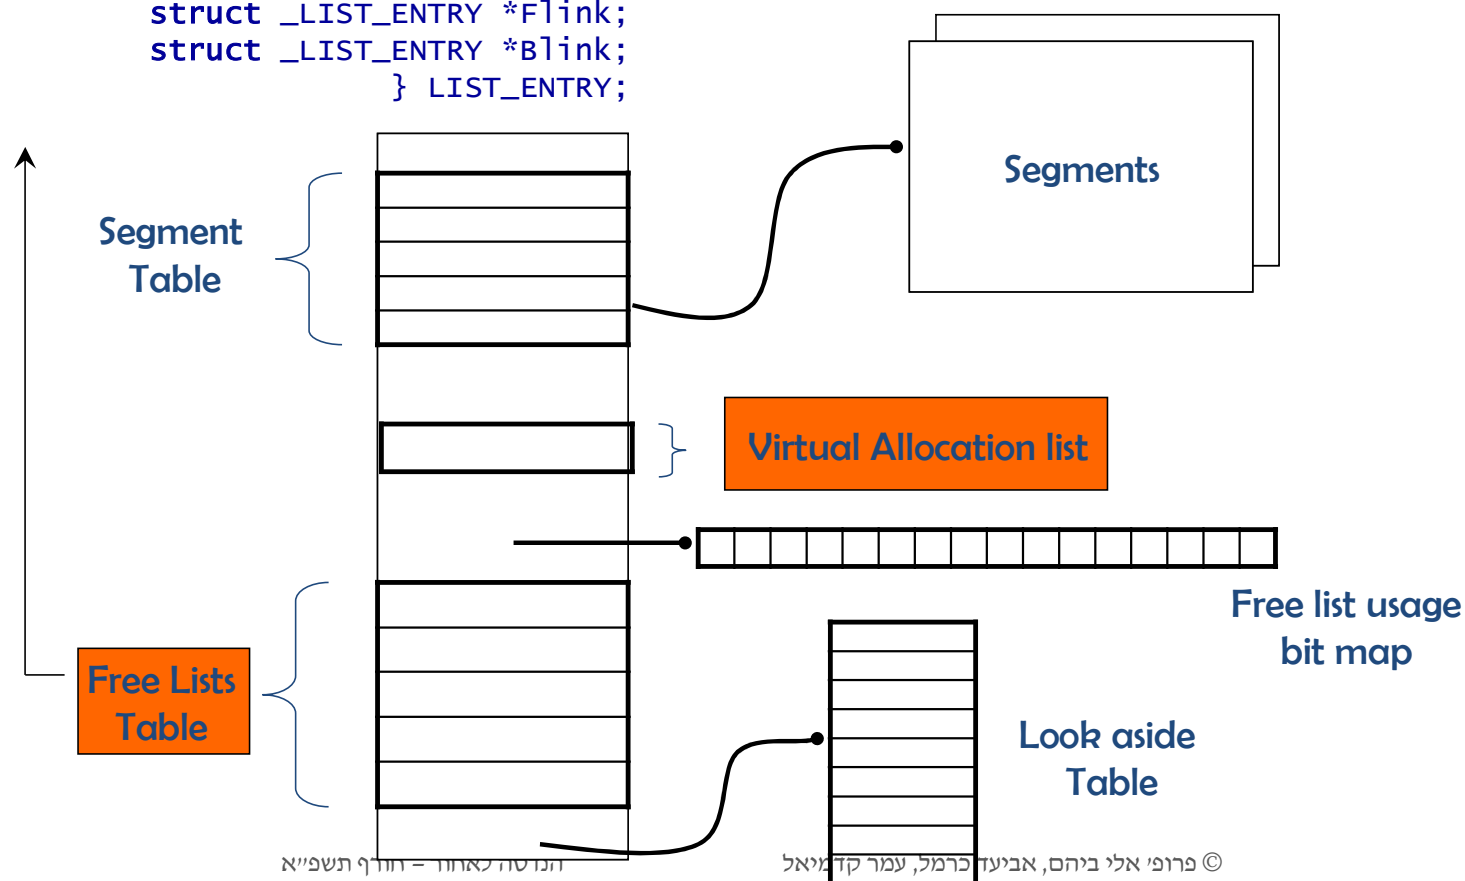
# The Heap API





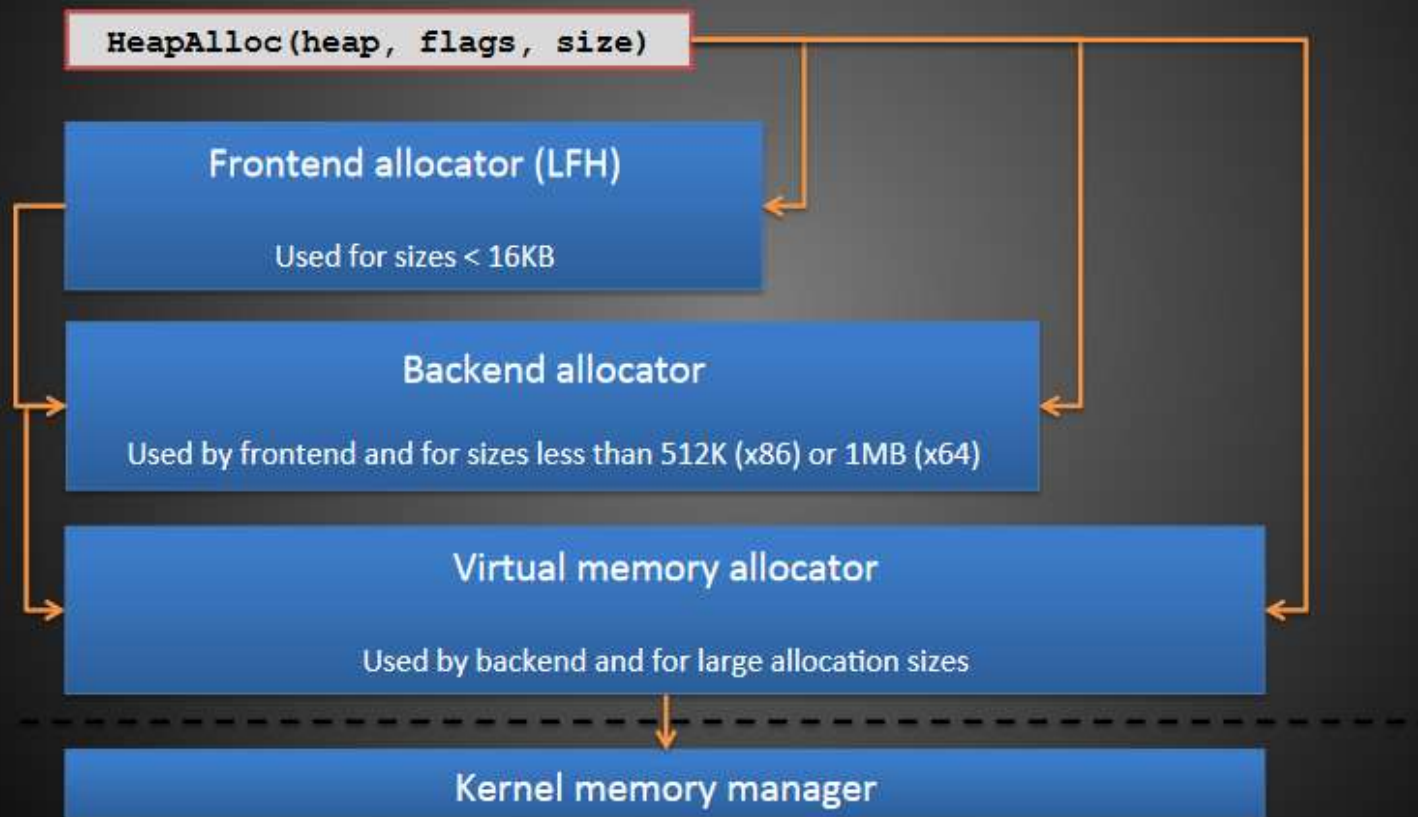
# The Heap structures

```
typedef struct _LIST_ENTRY {
    struct _LIST_ENTRY *Flink;
    struct _LIST_ENTRY *Blink;
} LIST_ENTRY;
```



# Windows 8 heap architecture

The general design of the Windows heap is unchanged in Windows 8



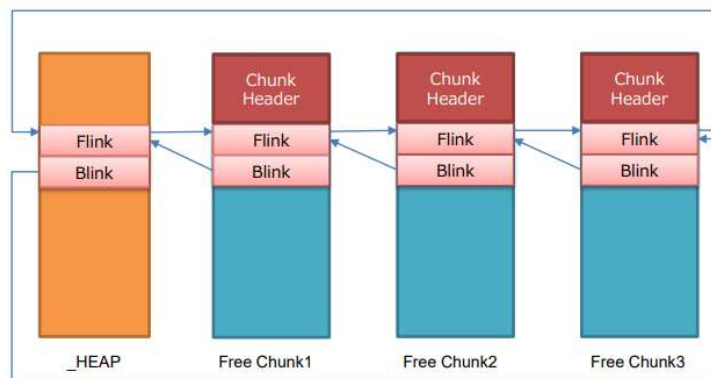
# Heap Structure

- Windows heap manager consist of mainly following 2 components.
  - Frontend
  - Backend
- Frontend is an interface to an application
  - Optimizes allocating small memory blocks
  - If it is able to respond to a request, it returns a memory block
  - If not, pass the request to the backend.
  - There are 2 frontend implementations.
    - Lookaside List (LAL) on Windows XP
    - Low Fragmentation Heap on Vista or later.



# Managing Free Chunks (Backend)

- As an application calling HeapAlloc or HeapFree in variety of order allocated chunks and free chunks are fragmented.
- To manage free chunks Windows heap manager uses doubly cyclic linked list.
- Free chunks also have a chunk header




※ In Windows XP actual free chunks are not managed as a single linked list but as a multiple linked lists. It is simplified here for the purpose of explanation of the exploit.  
© פרופ' אלי ביהם, אביעד כרמל, עמר קדמיאל



# The Heap structures

The heap management structures reside in the heap!



When a heap is first created there are two pointers that point to the first free block set in `FreeList[0]`. Assuming the heap base address is `0x00350000` then first available block can be found at `0x00350688`.

`0x00350178 (FreeList[0].Flink) = 0x00350688 (First Free Block)`  
`0x0035017C (FreeList[0].Blink) = 0x00350688 (First Free Block)`

`0x00350688 (First Free Block) = 0x00350178 (FreeList[0])`  
`0x0035068C (First Free Block+4) = 0x00350178 (FreeList[0])`




# Heap Structures

The instructions to remove an entry from a double linked list are:

```
prev_chunk->FLink = next_chunk  
next_chunk->BLink = prev_chunk
```

Where next\_chunk is this->FLink  
and prev\_chunk is this->BLink



# Heap smashing

If we assume that `prev_chunk->FLink` is loaded in ECX  
(and since `prev_chunk->FLink == &Prev_chunk + 0` in  
ECX is `prev_chunk` )  
And that `next_chunk` is in EAX...

If we build a fake header with `prev_chunk` and  
`next_chunk` of our choiche we have (from the previous  
code)

```
mov dword ptr [ecx],eax
```

EAX (Address A) written in the address  
pointed by ECX (Address B)

Arbitrary memory overwrite will happen on  
free of our faked chunk!





# The Heap structures

Access violation / Memory overwrite

Observe the following code:


```
mov dword ptr [ecx],eax
```

If we own both EAX and ECX we have an *arbitrary DWORD overwrite*. We can overwrite the data at any 32bit address with a 32bit value of our choosing.

In `RtlHeapFree` we <sup>had</sup> ~~have~~ such a line of code!!



# Heap smashing: Coalesce Algorithm



$B.Blink.Flink(A.Flink) = B.Flink(C)$   
 $B.Flink.Blink(C.Blink) = B.Blink(A)$

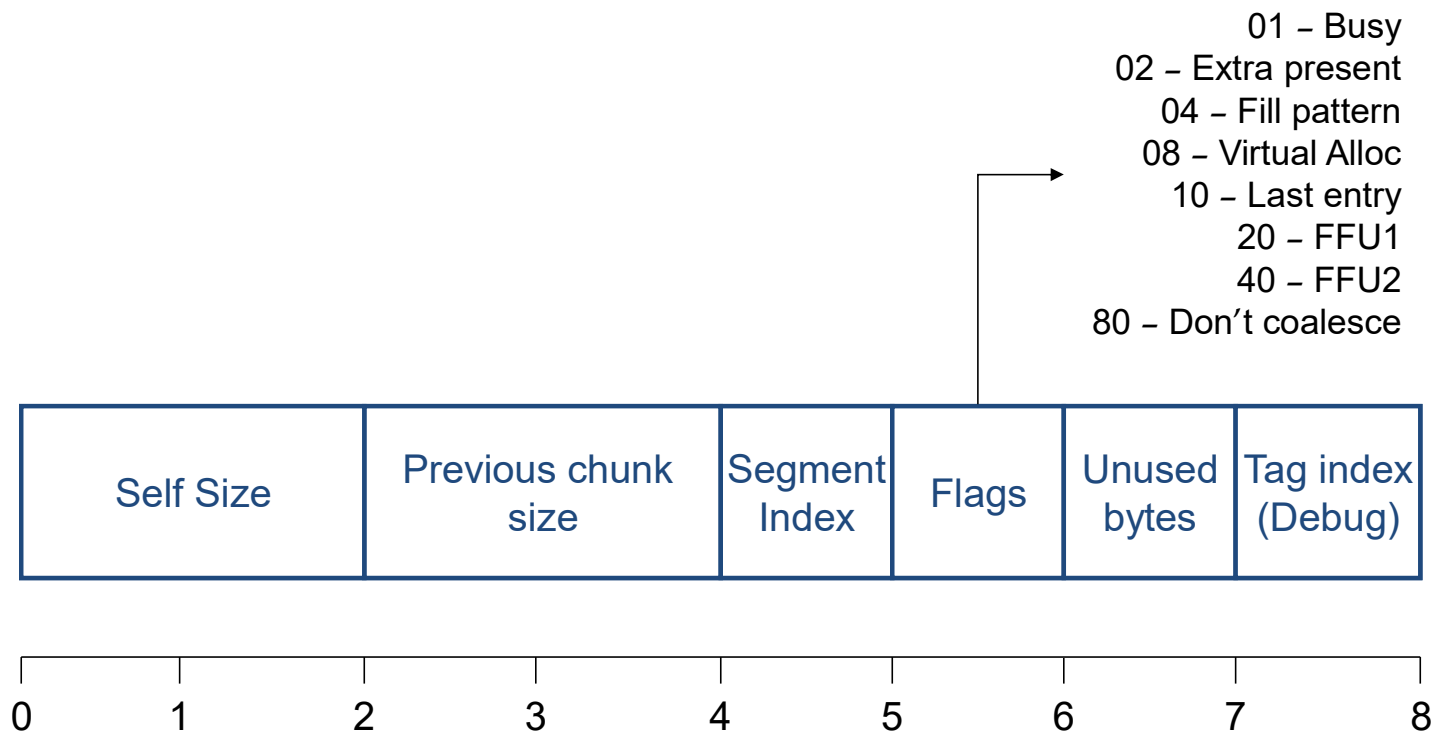


$*(B.Blink) = B.Flink$   
 $\Rightarrow *AddressB = AddressA$

**Arbitrary memory overwrite!**

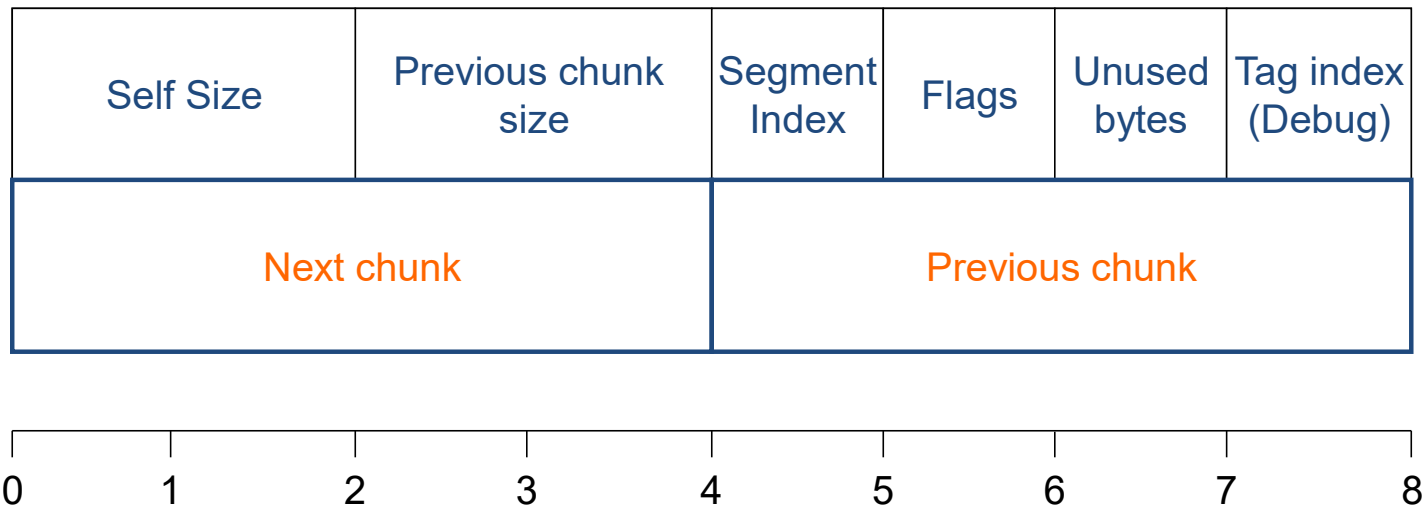


# The Heap structures



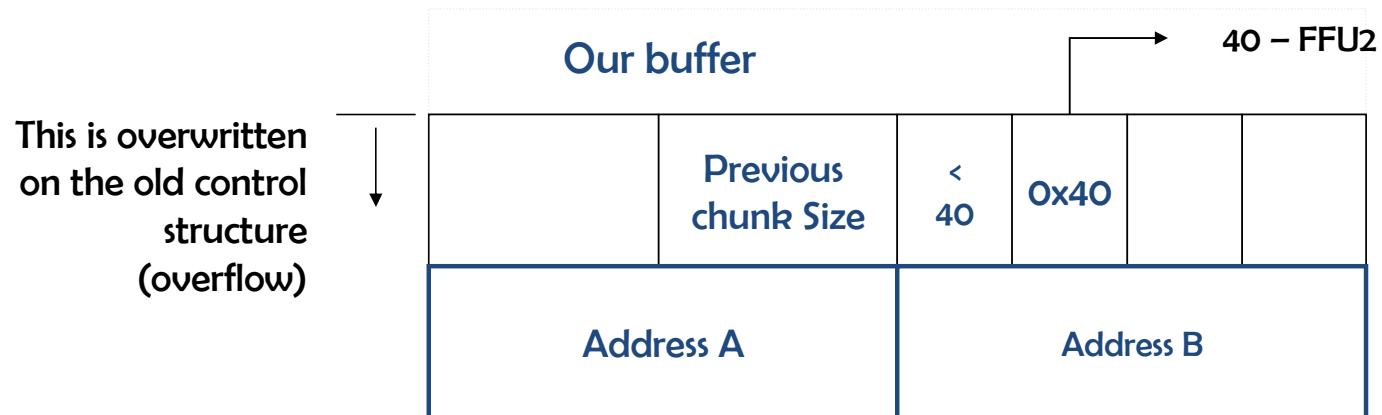
# The Heap structures

## Free chunk structure – 16 Bytes



# Heap Smashing - Overflow

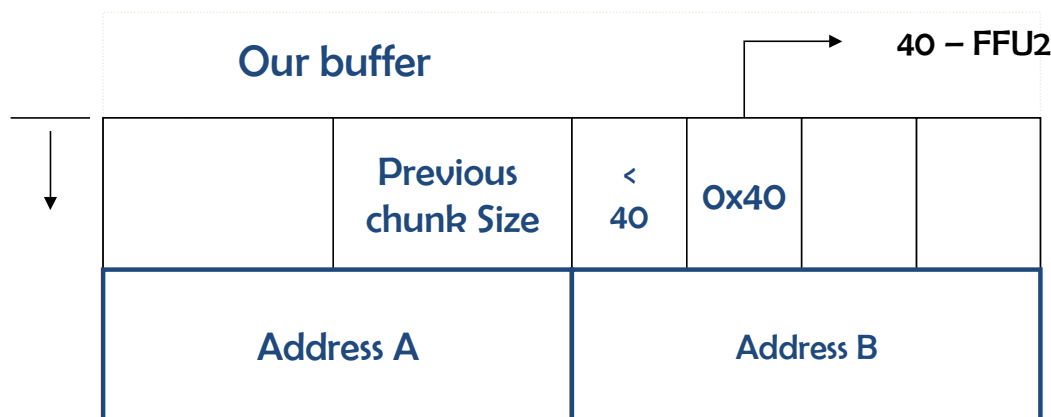
- Exploit: fake a freed buffer



# Heap Smashing – Double Free

- בוצע Free וה-chunk יצא מתוך הרשימה המקושרת.
  - כלומר ה-kernel ביצע unlink.

- מטרת התוקף עכשיו "להשתלט" על אותו קטע זיכרון ולייצר buffer שנראה נכון
  - כדי שכה-kernel ינסה לשחרר אותו בפעם הבאה – לא תהיה שגיאה (לפחות לפני שמה שהתוקף רוצה יתבצע)




# Heap Smashing – Double Free

• עכשיו קוראים ל-free עוד פעם

▪ ולכן מנסים להוציא את האיבר (שלא שם) מהרשימה

▪ ומבצעים

```
prev_chunk->FLink = next_chunk  
next_chunk->BLink = prev_chunk
```





# Heap Smashing – Double Free

- מה שמתבצע זה (לאחר חישוב):

Mov [v1+c], v2

- כאשר  $v1$  ו- $v2$  נמצאים בשליטה מלאה של התוקף ו-C ידוע  
▪ הנחת העבודה כאן היתה שמי שמבצע את התהליך הוא לא עוין.
- כלומר התוקף קיבל WWW בהרשאות Kernel.



# באינוקס זה לא היה קורה!



# Dynamic Memory Management in C

- Memory allocation: `malloc(size_t n)`
  - Allocates n bytes and returns a pointer to the allocated memory; memory not cleared
  - Also `calloc()`, `realloc()`
- Memory deallocation: `free(void * p)`
  - Frees the memory space pointed to by p, which must have been returned by a previous call to `malloc()`, `calloc()`, or `realloc()`
  - If `free(p)` has already been called before, undefined behavior occurs
  - If p is NULL, no operation is performed



# The Unlink Macro

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

Removes a chunk from a free list - when?





# Heap smashing: repairing the heap

- Many of the Windows API calls use the default process heap.
- After the overflow the heap is corrupt, so there will be surely an access violation.
- We then repair the heap following Litchfield's method: we reset the heap making it "appear" as if it is a fresh new heap.



# Heap smashing: repairing the heap

- Get a pointer to the Thread Information Block at fs:[18]
- Get a pointer to the Process Environment Block from the TEB.
- Get a pointer to the default process heap from the PEB
- We now have a pointer to the heap. Read the TotalFreeSize dword of the heap structure (at offset 0x28)
- Write this to our heap control structure.
- In the heap control structure, also set the flags to 0x14 (first segment) and the next 2 bytes to 0
- At heap\_base+0x178 we have FreeLists[0]. Set FreeLists[0].Flink and write edx into FreeLists[0].Blink to this
- Finally set the pointers at the end of our block to point to FreeLists[0]



## Arbitrary memory write

- So, as in the data pointer overwrite, we have an arbitrary DWORD memory overwrite
- This can be used to overwrite particular function pointers (Exception handlers VEH + SEH, atexit(), stack cookie exception handler, Lock and Win32k function pointers in PEB)





# Mitigations in Windows XP SP2

- Cookie in chunk header
  - 8bit checksum(cookie) is introduced in chunk header.
  - By validating its value it can detect overwrite of a chunk header.
  - The value of a cookie is based on the address of the chunk.
- Safe unlinking
  - Before removing an element from doubly linked list it checks if following condition is met:

$$[\text{Chunk}] \rightarrow \text{Flink} \rightarrow \text{Blink} == [\text{Chunk}] \rightarrow \text{Blink} \rightarrow \text{Flink} == [\text{Chunk}]$$

(Confirming if next/prev elements also point to the element)

- PEB Randomization
  - Randomize the address of PEB(Process Environment Block).
  - PEB contains values and addresses used by attackers not only in heap exploit.
  - This randomization makes the success rate of attacks low.



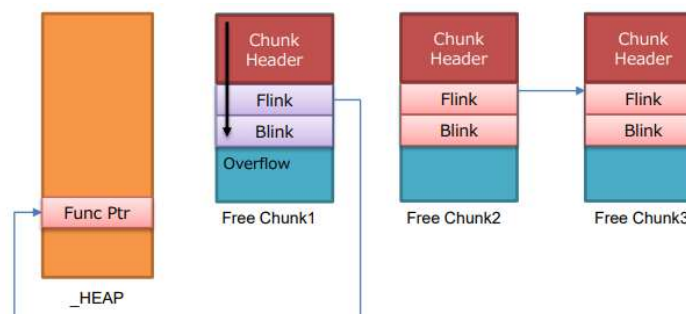
# Bypassing WinXP SP2 mitigations

- Using lookaside list
  - Freeing a chunk via HeapFree is handled by lookaside list first
  - Lookaside lists are singly linked lists of free chunks in each sizes.
  - If memory chunk is requested via HeapAlloc, it first checks if there is a free chunk in lookaside list and returns it if one exists.
- Important 2 mitigations does not work on lookaside list
  - Allocating a chunk from lookaside list does not make use of cookie
  - Safe Unlinking is not done (because it is not doubly linked list)



# Bypassing WinXP SP2 mitigations

- Using lookaside list
  - Overflow to make Flink have an address of a region containing a function pointer
  - If data written in the region allocated from subsequent second HeapAlloc can be controlled, the pointer can be rewritten by an arbitrary address.
  - In the figure below, Flink uses a function pointer in \_HEAP structure (This function pointer is used and called in heap management process)



# Windows Vista

- Low Fragmentation Heap
  - Lookaside list is replaced with Low Fragmentation Heap
  - Impossible to attack using lookaside list
- Randomizing Block Metadata
  - Chunk header is xored with `_HEAP->Encoding`
  - Overwriting chunk header with predicted cookie still results in unexpected state of the chunk header.
- Enhanced entry header cookie
  - Cookie value also checks values in chunk header
  - Cookie had been calculated based on the chunk address but now it is calculated/validated with values in a chunk header.
- Heap base randomization
  - The base address of `_HEAP` structure is randomized – Makes it difficult to overwrite data in `_HEAP` structure
- Heap function pointer encoding
  - A function pointer in `_HEAP` structure is xored with a value
  - Mitigations for rewriting a function pointer in `_HEAP` structure



# LFH design changes & integrity checks


Change in Windows 8	Impact
LFH is now a bitmap-based allocator	LinkOffset corruption no longer possible [8]
Multiple catch-all EH blocks removed	Exceptions are no longer swallowed
HEAP handle can no longer be freed	Prevents attacks that try to corrupt HEAP handle state [7]
HEAP CommitRoutine encoded with global key	Prevents attacks that enable reliable control of the CommitRoutine pointer [7]
Validation of extended block header	Prevents unintended free of in-use heap blocks [7]
Busy blocks cannot be allocated	Prevents various attacks that reallocate an in-use block [8,11]
Heap encoding is now enabled in kernel mode	Better protection of heap entry headers [19]

**Outcome: attacking metadata used by the heap is now even more difficult**



# Guard pages

- Guard pages are now used to partition the heap
  - Designed to prevent & localize corruption in some cases
  - Touching a guard page results in an exception



- Insertion points for guard pages are constrained
  - Large allocations
  - Heap segments
  - Max-sized LFH subsegments (probabilistic on 32-bit)



# Allocation order randomization

- Allocation order is now nondeterministic (LFH only)
  - Exploits often rely on surgical heap layout manipulation [10]
  - Randomization makes heap normalization unreliable

Windows 7 LFH block allocation behavior



Windows 8 LFH block allocation behavior



- Maximizing reliability is more challenging
  - Application-specific and vulnerability-specific
  - May require corrupting more data (increasing instability)
  - May require allocating more data (triggering guard pages)

# לסיכום..

- ניצול חולשות Heap קשה ומסובך הרבה יותר מחולשות Stack
  - וגם תלוי מימוש וארכיטקטורה – דורש RE
- התקדמות משמעותית גם במנגנוני ההגנה ב-Heap
  - בעיקר רנדומיזציה ו-Guard Pages
- אבל ב-Heap קיימים מנגנונים מסובכים הרבה יותר
  - שלפיכך פגיעים יותר לשגיאות

